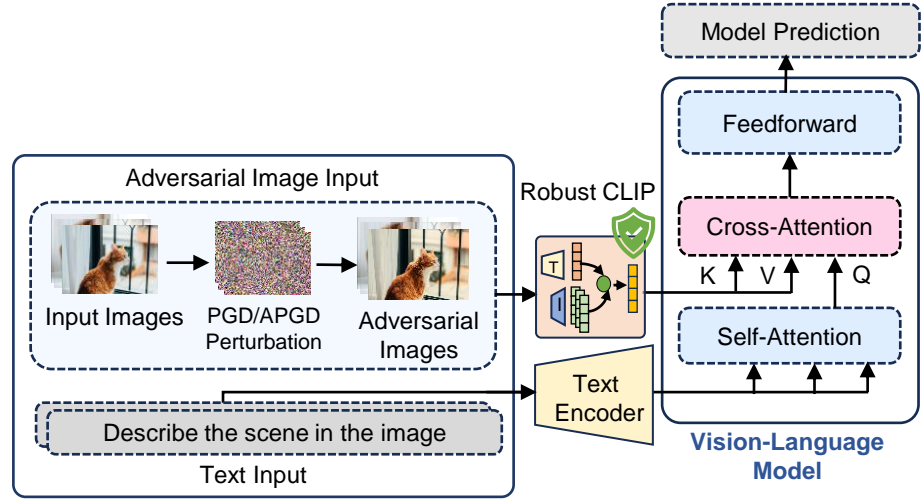


(a) Unsupervised adversarial fine-tuning



(b) VLM Inference